

Assignment #1: Building an OS Shell

Due: February 2, 2026 at 23:59

1. Infrastructure Description

Welcome to the first OS assignment where we will build an OS Shell!

This is the first of a series of three assignments that build upon each other. By the end of the semester, you will have a running simulation of a simple Operating System and you will get a good idea of how the basic OS modules work.

You will use the **C programming language** in your implementation, since most of the practical operating systems kernels are written in the C/C++ (e.g., including Windows, Linux, MacOS, and many others).

The assignment is presented from a Linux point of view. We will run the assignments in the Ubuntu 24.04.3 (Noble Numbat) release. For ease of use, we will install Ubuntu as a virtual machine, via VirtualBox (see the VirtualBox tutorial on MyCourses).

The TAs will also test the assignments in the VirtualBox + Ubuntu 24.04.3 environment.

TA responsible for grading and Q&A for this assignment: Siam Antar siam.antar@mail.mcgill.ca

TA responsible for infrastructure troubleshooting: Olivier Michaud olivier.michaud2@mail.mcgill.ca

1.1 Starter files description:

We provided you with a simple shell to start with, which you will enhance in this assignment. Once you sync your fork with ours, the starter code will be in [A1/starter-code](#). Whether or not you want to use the starter code, run `mkdir -p project/src` to create the src folder where you will write your code. If you want to use the starter code, now run

```
cp A1/starter-code/* project/src/
```

Navigate to the src folder with cd and take a moment to get familiar with the code.

- Use the following command to compile: `make mysh`
- Re-compiling your shell after making modifications: `make clean; make mysh`

Running your starter shell

- **Interactive mode:** From the command line prompt type: `./mysh`
- **Batch mode:** You can also use input files to run your shell. To use an input file, from the command line prompt type: `./mysh < testfile.txt`

Starter shell interface. The starter shell supports the following commands:

COMMAND	DESCRIPTION
<code>help</code>	<i>Displays all the commands</i>
<code>quit</code>	<i>Exits / terminates the shell with “Bye!”</i>
<code>set VAR STRING</code>	<i>Assigns a value to shell memory</i>
<code>print VAR</code>	<i>Displays the STRING assigned to VAR</i>
<code>source SCRIPT.TXT</code>	<i>Executes the file SCRIPT.TXT</i>

More details on command behavior:

- The commands are case sensitive.
- If the user inputs an unsupported command the shell displays “Unknown command”.
- `set VAR STRING` first checks to see if VAR already exists. If it does exist, STRING overwrites the previous value assigned to VAR. If VAR does not exist, then a new entry is added to the shell memory where the variable name is VAR and the contents of the variable is STRING. For now, each value assigned to a variable is a single alphanumeric token (i.e., no special characters, no spaces, etc.). For example:
 - `set x 10` creates a new variable x and assigns to it the string 10.
 - `set name Bob` creates a new variable called name with string value Bob.
 - `set x Mary`, replaced the value 10 with Mary.
- `print VAR` first checks to see if VAR exists. If it does not exist, then it displays the error “Variable does not exist”. If VAR does exist, then it displays the STRING. For example: `print x` from the above example will display Mary.
- `source SCRIPT.TXT` assumes that a text file exists with the provided file name, *in the current directory*. It opens that text file and then sends each line one at a time to the interpreter. The interpreter treats each line of text as a command. At the end of the script, the file is closed, and the command line prompt is displayed once more. While the script executes the command line prompt is not displayed. If an error occurs while executing the script due a command syntax error, then the error is displayed, and the script continues executing.

1.2 Your tasks:

Your task is to add the following functionality to the starter shell.

1.2.1. Add the `echo` command.

The `echo` command is used for displaying strings which are passed as arguments on the command line. This simple version of `echo` only takes **one token string** as input. The token can be:

- **An alphanumeric string.** In this case, `echo` simply displays the string on a new line and then returns the command prompt to the user.

Example execution (interactive mode):

```
$ echo mary
mary
$
```

- **An alphanumeric string preceded by \$.** In this case, `echo` checks the shell memory for a variable that has the name of the alphanumeric string following the \$ symbol.
 - If the variable is found, `echo` displays the value associated to that variable, similar to the `print` command and then returns the command prompt to the user.
 - If the variable is not found, `echo` displays an empty line and then returns the command prompt to the user.

Example execution (interactive mode):

```
$ echo $mary
// blank line
$
$ set mary 123
$ echo $mary
123
$
```

Assumptions:

- You can assume that the token string is <100 characters.

1.2.2. Enhance batch mode execution.

1. Batch mode execution in your starter shell enters an infinite loop if the last command in the input file is not `quit`. Fix this issue so the shell does not enter an infinite loop. Instead, the shell should terminate after running all the instructions in the input file.
2. Batch mode execution in your starter shell displays \$ for every line of command in the batch mode. Change the batch execution so that \$ is only displayed in the interactive mode.

This step is extremely important – if you do not complete this step, you will fail every test.

1.2.3. Add the ls, mkdir, touch, and cd commands.

Add three new commands to your shell:

1. `my_ls` lists all the files present in the *current directory*.
 - If the current directory contains other directories, `my_ls` displays only the name (not the contents) of the directory.
 - Each file or directory name needs to be displayed on a separate line.
 - The file/directory names are shown in alphabetical order, similar to the `sort` command in Linux:

- Names starting with a number will appear before lines starting with a letter.
 - Names starting with a letter that appears earlier in the alphabet will appear before lines starting with a letter that appears later in the alphabet.
 - Names starting with an uppercase letter will appear before lines starting with the same letter in lowercase.
- Additionally, `my_ls` should include the special entries `."` and `..` in the output, which represent the current directory and the parent directory, respectively.
 - Hint: investigate the `dirent.h` header file.
2. `my_mkdir dirname` creates a new directory with the name `dirname` in the *current directory*.
- `dirname` can be (1) an alphanumeric string, or (2) an alphanumeric string preceded by `$`.
 - If `dirname` is an alphanumeric string, `my_mkdir` creates a directory with the given name.
 - If `dirname` is an alphanumeric string preceded by `$`, `my_mkdir` checks the shell memory for a variable that has the name of the alphanumeric string following the `$` symbol.
 - If the variable exists in the shell memory and contains a single alphanumeric token, `my_mkdir` creates a directory using the value associated to that variable as the directory name.
 - If the variable is not found or the variable contains something other than a single alphanumeric token, `my_mkdir` displays *“Bad command: my_mkdir”* and then returns the command prompt to the user.
3. `my_touch filename` creates a new empty file inside the current directory. `filename` is an alphanumeric string.
4. `my_cd dirname` changes current directory to directory `dirname`, inside the current directory. If `dirname` does not exist inside the current directory, `my_cd` displays *“Bad command: my_cd”* and stays inside the current directory. `dirname` should be an alphanumeric string, you do not need to consider the case where `dirname` is a shell variable.

Assumptions:

- You can assume that file/directory names are <100 characters.

1.2.4. One-liners.

The starter shell only supports a single command per line. This is not the case for regular shells where multiple commands can be chained. Your task is to implement a simple chaining of instructions, where the shell can take as input multiple commands separated by semicolons (the ; symbol).

Assumptions:

- The instructions separated by semicolons are executed one after the other.
- The total length of the combined instructions does not exceed 1000 characters.
- There will be at most 10 chained instructions.
- Semicolon is the only accepted separator. (Spaces are still ignored, of course.)

Example execution (interactive mode):

```
$ set x abc; set y 123; print y; print x
123
abc
$
```

1.2.5. Implementing the 'run' Command with Fork-Exec-Wait.

Add the command `run` which uses the “fork-exec-wait” pattern discussed in class to invoke other commands. That is, it forks the shell and calls one of the flavors of exec (see the man page) to execute the given command. You should use a flavour of exec that lets you pass the rest of the user input as command-line arguments to the command.

You will need to read the documentation for exec and wait. You can get the man pages from your terminal with `man 2 wait` and `man 3 exec`, or you can look it up online.

Example:

```
$ run cat shell.h
#define MAX_USER_INPUT 1000
int parseInput(char inp[]);
$
```

Note: The ‘`run`’ command is only needed to ensure that our testcases work properly, in particular that the Bad Command error still appears when it should. For fun (and in a different branch so you don’t accidentally break your submission), you can extend your shell into a proper shell by removing the `run` command and attempting to fork-exec *any* command name that isn’t built into the shell. You could also try adding more chaining operators like `&&` and `||`, or you can try adding file redirection or pipes. The world is your oyster!

2. TESTCASES

We provide you with 10 testcases and expected outputs for your code in the starter code repository. Please run the testcases to ensure your code runs as expected, and make sure you get similar results in the automatic tests. These are all the test cases that the TAs will run when evaluating your work.

To run a test, navigate to `A1/test-cases`. Then run `../../project/src/mysh < [testfile].txt`. Alternatively, you can move or copy the mysh executable to the test directory to shorten the path. Use `diff` to compare the output with the expected output. **We highly recommend writing a bash or python script to help you run your tests automatically.** Writing testing scripts is part of C programming!

IMPORTANT: You can assume that the TAs will run one test at a time **in batch mode**. Regrade requests of the form “I never tested my code in batch mode, but it works in interactive mode” will be denied.

3. WHAT TO HAND IN

The assignment is **due on Feb 2, 2026 at 23:59, on MyCourses, in the Assignments tab.**

In addition to the code, please replace the README at the top level of the repository. Your README should include any additional comments the author(s) would like the TA to see, and mention whether the code uses the starter code provided by the OS team or not.

The project must compile by running `make clean; make mysh`.

The project must run in batch mode, i.e. `./mysh < testfile.txt`

4. HOW IT WILL BE GRADED

Your program must compile and run to be graded. If the code does not compile/run using the commands in Section 3, you will receive **0 points** for the entire assignment. If you think your code is correct and there is an issue with the VM setup, contact TA Olivier Michaud at olivier.michaud2@mail.mcgill.ca.

We will use **your Makefile**. You may change your Makefile however you wish to fit your project.

Your assignment is graded out of 100 points. Up to 80 of these points are awarded for passing test cases. The other 20 are awarded by a TA performing code review.

Test Cases: You were provided 10 test cases, with expected outputs. If your code matches the expected output, you will receive 8 points for each testcase. You will receive 0 points for each test case where your output does not match the expected output. When comparing output, we ignore differences in capitalization and whitespace. Your output must have the same words, in the same order, as the expected output.

Any test cases for which you have hardcoded the output will receive 0 points. “Hardcoding” means you’ve copied the expected output directly to your program output, or otherwise obtained the correct output without implementing the assignment requirements yourself. Note that this means that you cannot use the `system` function that you saw in COMP206, nor other functions with similar behavior like `popen`. Helper functions in the C standard library, such as those in `<dirent.h>`, are allowed.

Code Review: If your code runs at all, regardless of how many points it scores (even 0!), a TA will perform code review. This entails looking at your code, spending some time trying to understand it, and then evaluating it for how straightforward it was to understand and for systems design and structure.

- **Your code must have useful, meaningful comments.** The TA will be trying to understand your code, and they didn’t write it! Part of writing systems software is being able to write an expansive program that *other people can understand*. Whenever your code is doing something non-trivial, you should probably have a comment explaining why.
- **Your software system should be reasonably organized and structured.** A helper function that is used in only one place should be near the function that uses it. A function that’s part of an API, e.g. the `shellmemory`, should be in the file implementing that API. If you need to design a new API, such as for a new data structure like a queue or PCB, that API should get its own header/C file pair. Code related to parsing the input and running the shell’s outer interactive loop probably belongs in `shell`. Code implementing the behavior of a shell command probably belongs in `interpreter`. In general, don’t be afraid to add new files to the project. Having more header files with APIs open on the side, and trusting that their APIs are implemented

correctly, is a low mental burden. Having more C code open on the side, and having to read the implementations to see how they work, is a high mental burden. We want low burdens 😊.

- **You cannot hardcode test cases.** See the warning under the ‘Test cases’ section. If the TA sees hardcoded, or the use of functions like `system` and `popen`, we will deduct all of the points from the affected test cases.
- **Your code must be written in C.** If the TA goes to review your code, and finds a project written in a language other than C (even if you’ve correctly set up a Makefile for that language....) you will get a zero. **C++ is not C.**